

This notebook contains a generator class for Keras called `BSONIterator` that can read directly from the BSON data. You can use it in combination with `ImageDataGenerator` for doing data augmentation.

▼ start

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!nvidia-smi
```

```
Fri Sep 9 06:00:08 2022
```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		460.32.03		Driver Version: 460.32.03			CUDA Version: 11.2		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
0	Tesla	P100-PCIE...	Off	00000000:00:04.0		Off	0		
N/A	41C	P0	27W / 250W	0MiB / 16280MiB		0%	Default		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
=====									
No running processes found									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

```
! pip install tensorflow==1.15.5
! pip install keras==2.0.8
! pip install tensorflow-gpu==1.15.5
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public
Requirement already satisfied: tensorflow==1.15.5 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: keras-applications>=1.0.8 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: numpy<1.19.0,>=1.16.0 in /usr/local/lib/python3.7/dist-packages
```

```

Requirement already satisfied: tensorboard<1.16.0,>=1.15.0 in /usr/local/lib/python3.7/c
Requirement already satisfied: tensorflow-estimator==1.15.1 in /usr/local/lib/python3.7/
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.7/di
Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.7/dist-pack
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: h5py<=2.10.0 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: gast==0.2.2 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packa
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: typing-extensions>=3.6.4 in /usr/local/lib/python3.7/dist
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Requirement already satisfied: keras==2.0.8 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: pyyaml in /usr/local/lib/python3.7/dist-packages (from ke
Requirement already satisfied: scipy>=0.14 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: numpy>=1.9.1 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.7/dist-packages (fro
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Requirement already satisfied: tensorflow-gpu==1.15.5 in /usr/local/lib/python3.7/dist-p
Requirement already satisfied: gast==0.2.2 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: protobuf>=3.6.1 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: keras-applications>=1.0.8 in /usr/local/lib/python3.7/di
Requirement already satisfied: numpy<1.19.0,>=1.16.0 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: absl-py>=0.7.0 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: astor>=0.6.0 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: tensorboard<1.16.0,>=1.15.0 in /usr/local/lib/python3.7/c
Requirement already satisfied: google-pasta>=0.1.6 in /usr/local/lib/python3.7/dist-pack
Requirement already satisfied: grpcio>=1.8.6 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: keras-preprocessing>=1.0.5 in /usr/local/lib/python3.7/di
Requirement already satisfied: tensorflow-estimator==1.15.1 in /usr/local/lib/python3.7/
Requirement already satisfied: wrapt>=1.11.1 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: h5py<=2.10.0 in /usr/local/lib/python3.7/dist-packages (f
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.7/dist-packa
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7/dist-packag
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: typing-extensions>=3.6.4 in /usr/local/lib/python3.7/dist

```

```

!pip install kaggle
from google.colab import files
from datetime import datetime
api_token = files.upload()

```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public>
 Requirement already satisfied: kaggle in /usr/local/lib/python3.7/dist-packages (1.5.12)
 Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: python-slugify in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: python-dateutil in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: certifi in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from kaggle)
 Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from kaggle)

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable

```
!mkdir ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

```
!kaggle competitions download -c 'cdiscout-image-classification-challenge' -f train.bson
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you
 Downloading train.bson.zip to /content
 100% 47.2G/47.2G [05:17<00:00, 151MB/s]
 100% 47.2G/47.2G [05:17<00:00, 160MB/s]



```
!kaggle competitions download -c 'cdiscout-image-classification-challenge' -f category_names
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you
 Downloading category_names.csv to /content
 0% 0.00/406k [00:00<?, ?B/s]
 100% 406k/406k [00:00<00:00, 102MB/s]



```
!unzip /content/train.bson.zip
```

Archive: /content/train.bson.zip
 inflating: train.bson

```
!rm /content/train.bson.zip
```

```
import os, sys, math, io
import numpy as np
import pandas as pd
import multiprocessing as mp
import bson
import struct

%matplotlib inline
import matplotlib.pyplot as plt
```

```

import keras
from keras.preprocessing.image import load_img, img_to_array
import tensorflow as tf

from collections import defaultdict
from tqdm import *

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list the files in

from subprocess import check_output
print(check_output(["ls", "/content"]).decode("utf8"))

# Any results you write to the current directory are saved as output.

    Using TensorFlow backend.
    category_names.csv
    drive
    kaggle.json
    sample_data
    train.bson

keras.__version__, tf.__version__

('2.0.8', '1.15.5')

import tensorflow as tf
print(tf.__version__, tf.test.is_gpu_available())

1.15.5 True

data_dir = "/content"

train_bson_path = os.path.join(data_dir, "train.bson")
num_train_products = 7069896

# train_bson_path = os.path.join(data_dir, "train_example.bson")
# num_train_products = 82

test_bson_path = os.path.join(data_dir, "test.bson")
num_test_products = 1768182

```

▼ Part 1: Create lookup tables

The generator uses several lookup tables that describe the layout of the BSON file, which products and images are part of the training/validation sets, and so on.

You only need to generate these tables once, as they get saved to CSV files. If you already have these CSV files, skip to part 2.

▼ Lookup table for categories

```
categories_path = os.path.join(data_dir, "category_names.csv")
categories_df = pd.read_csv(categories_path, index_col="category_id")

# Maps the category_id to an integer index. This is what we'll use to
# one-hot encode the labels.
categories_df["category_idx"] = pd.Series(range(len(categories_df)), index=categories_df.index)

categories_df.to_csv("categories.csv")
categories_df.head()
```

category_id	category_level1	category_level2	category_level3	category_idx
1000021794	ABONNEMENT / SERVICES	CARTE PREPAYEE	CARTE PREPAYEE MULTIMEDIA	0
1000012764	AMENAGEMENT URBAIN - VOIRIE	AMENAGEMENT URBAIN	ABRI FUMEUR	1
1000012776	AMENAGEMENT URBAIN - VOIRIE	AMENAGEMENT URBAIN	ABRI VELO - ABRI MOTO	2
1000012788	AMENAGEMENT	AMENAGEMENT	CONTAINES A PAILLON	3

Create dictionaries for quick lookup of category_id to category_idx mapping.

```
def make_category_tables():
    cat2idx = {}
    idx2cat = {}
    for ir in categories_df.iterrows():
        category_id = ir[0]
        category_idx = ir[4]
        cat2idx[category_id] = category_idx
        idx2cat[category_idx] = category_id
    return cat2idx, idx2cat
```

```
cat2idx, idx2cat = make_category_tables()
```

Show hidden output

```
# Test if it works:
```

```
cat2idx[1000012755], idx2cat[4]
```

▼ Read the BSON files

We store the offsets and lengths of all items, allowing us random access to the items later.

Inspired by code from: <https://www.kaggle.com/vfdev5/random-item-access>

Note: this takes a few minutes to execute, but we only have to do it once (we'll save the table to a CSV file afterwards).

```
def read_bson(bson_path, num_records, with_categories):
    rows = {}
    with open(bson_path, "rb") as f, tqdm(total=num_records) as pbar:
        offset = 0
        while True:
            item_length_bytes = f.read(4)
            if len(item_length_bytes) == 0:
                break

            length = struct.unpack("<i", item_length_bytes)[0]

            f.seek(offset)
            item_data = f.read(length)
            assert len(item_data) == length

            item = bson.BSON.decode(item_data)
            product_id = item["_id"]
            num_imgs = len(item["imgs"])

            row = [num_imgs, offset, length]
            if with_categories:
                row += [item["category_id"]]
            rows[product_id] = row

            offset += length
            f.seek(offset)
            pbar.update()

    columns = ["num_imgs", "offset", "length"]
    if with_categories:
        columns += ["category_id"]

    df = pd.DataFrame.from_dict(rows, orient="index")
    df.index.name = "product_id"
    df.columns = columns
    df.sort_index(inplace=True)
    return df
```

```
%time train_offsets_df = read_bson(train_bson_path, num_records=num_train_products, with_cate
```

```
100%|██████████| 7069896/7069896 [10:21<00:00, 11380.28it/s]
CPU times: user 1min 28s, sys: 25.7 s, total: 1min 53s
Wall time: 10min 30s
```

```
train_offsets_df.to_csv("/content/drive/MyDrive/kerasgenmodel/meta_datatrain_offsets.csv")
```

```
#train_offsets_df.to_csv("/content/drive/MyDrive/kerasgenmodel/meta_datatrain_offsets.csv")
train_offsets_df = pd.read_csv('/content/drive/MyDrive/csvfiles/train_offsets.csv')
```

```
train_offsets_df.head()
```

	product_id	num_imgs	offset	length	category_id
0	0	1	0	6979	1000010653
1	1	1	6979	7318	1000010653
2	2	1	14297	5455	1000004079
3	3	1	19752	4580	1000004141
4	4	1	24332	6346	1000015539

```
#train_offsets_df.to_csv("/content/drive/MyDrive/csvfiles/train_offsets.csv")
```

```
# How many products?
len(train_offsets_df)
```

```
7069896
```

```
# How many categories?
len(train_offsets_df["category_id"].unique())
```

```
5270
```

```
# How many images in total?
train_offsets_df["num_imgs"].sum()
```

```
12371293
```

▼ Create a random train/validation split

We split on products, not on individual images. Since some of the categories only have a few products, we do the split separately for each category.

This creates two new tables, one for the training images and one for the validation images. There is a row for every single image, so if a product has more than one image it occurs more than once in the table.

```
def make_val_set(df, split_percentage=0.2, drop_percentage=0.):
    # Find the product_ids for each category.
    category_dict = defaultdict(list)
    for ir in tqdm(df.itertuples()):
        category_dict[ir[4]].append(ir[0])

    train_list = []
    val_list = []
    with tqdm(total=len(df)) as pbar:
        for category_id, product_ids in category_dict.items():
            category_idx = cat2idx[category_id]

            # Randomly remove products to make the dataset smaller.
            keep_size = int(len(product_ids) * (1. - drop_percentage))
            if keep_size < len(product_ids):
                product_ids = np.random.choice(product_ids, keep_size, replace=False)

            # Randomly choose the products that become part of the validation set.
            val_size = int(len(product_ids) * split_percentage)
            if val_size > 0:
                val_ids = np.random.choice(product_ids, val_size, replace=False)
            else:
                val_ids = []

            # Create a new row for each image.
            for product_id in product_ids:
                row = [product_id, category_idx]
                for img_idx in range(df.loc[product_id, "num_imgs"]):
                    if product_id in val_ids:
                        val_list.append(row + [img_idx])
                    else:
                        train_list.append(row + [img_idx])
                pbar.update()

    columns = ["product_id", "category_idx", "img_idx"]
    train_df = pd.DataFrame(train_list, columns=columns)
    val_df = pd.DataFrame(val_list, columns=columns)
    return train_df, val_df
```

Create a 80/20 split. Also drop 90% of all products to make the dataset more manageable. (Note: if `drop_percentage > 0`, the progress bar doesn't go all the way.)


```
train_images_df, val_images_df = make_val_set(train_offsets_df, split_percentage=0.1,
                                              drop_percentage=0.0)
```

```
train_images_df.head()
```

	product_id	category_idx	img_idx
0	0	5055	0
1	11	5055	0
2	16	5055	0
3	28	5055	0
4	30	5055	0

```
val_images_df.head()
```

```
print("Number of training images:", len(train_images_df))
print("Number of validation images:", len(val_images_df))
print("Total images:", len(train_images_df) + len(val_images_df))
```

```
Number of training images: 11139021
Number of validation images: 1232272
Total images: 12371293
```

Are all categories represented in the train/val split? (Note: if the drop percentage is high, then very small categories won't have enough products left to make it into the validation set.)

```
len(train_images_df["category_idx"].unique()), len(val_images_df["category_idx"].unique())

(5270, 5270)
```

Quickly verify that the split really is approximately 80-20:

```
category_idx = 619
num_train = np.sum(train_images_df["category_idx"] == category_idx)
num_val = np.sum(val_images_df["category_idx"] == category_idx)
num_val / num_train

0.11453077699293643
```

Close enough. ;-)

Remember that we split on products but not all products have the same number of images, which is where the slightly discrepancy comes from. (Also, there tend to be fewer validation images if `drop_percentage > 0.`)

Save the lookup tables as CSV so that we don't need to repeat the above procedure again.

```
train_images_df.to_csv("train_images.csv")
val_images_df.to_csv("val_images.csv")
```

```
train_images_df.to_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/train_images.csv")
val_images_df.to_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/val_images.csv")
```

▼ Part 2: The generator

First load the lookup tables from the CSV files (you don't need to do this if you just did all the steps from part 1).

```
categories_df = pd.read_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/categories.csv",
cat2idx, idx2cat = make_category_tables())
```

```
train_offsets_df = pd.read_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/train_offsets.c
train_images_df = pd.read_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/train_images.cs
val_images_df = pd.read_csv("/content/drive/MyDrive/kerasgenmodel/meta_data/val_images.csv",
```

Show hidden output

The Keras generator is implemented by the `BSONIterator` class. It creates batches of images (and their one-hot encoded labels) directly from the BSON file. It can be used with multiple workers.

Note: For fastest results, put the train.bson and test.bson files on a fast drive (SSD).

See also the code in: <https://github.com/fchollet/keras/blob/master/keras/preprocessing/image.py>

```
from keras.preprocessing.image import Iterator
from keras.preprocessing.image import ImageDataGenerator
from keras import backend as K
```

```
class BSONIterator(Iterator):
    def __init__(self, bson_file, images_df, offsets_df, num_class,
                 image_data_generator, lock, target_size=(180, 180),
                 with_labels=True, batch_size=32, shuffle=False, seed=None):

        self.file = bson_file
        self.images_df = images_df
        self.offsets_df = offsets_df
        self.with_labels = with_labels
        self.samples = len(images_df)
```

```

self.num_class = num_class
self.image_data_generator = image_data_generator
self.target_size = tuple(target_size)
self.image_shape = self.target_size + (3,)

print("Found %d images belonging to %d classes." % (self.samples, self.num_class))

super(BSONIterator, self).__init__(self.samples, batch_size, shuffle, seed)
self.lock = lock

def _get_batches_of_transformed_samples(self, index_array):
    batch_x = np.zeros((len(index_array),) + self.image_shape, dtype=K.floatx())
    if self.with_labels:
        batch_y = np.zeros((len(batch_x), self.num_class), dtype=K.floatx())

    for i, j in enumerate(index_array):
        # Protect file and dataframe access with a lock.
        with self.lock:
            image_row = self.images_df.iloc[j]
            product_id = image_row["product_id"]
            offset_row = self.offsets_df.loc[product_id]

            # Read this product's data from the BSON file.
            self.file.seek(offset_row["offset"])
            item_data = self.file.read(offset_row["length"])

            # Grab the image from the product.
            item = bson.BSON.decode(item_data)
            img_idx = image_row["img_idx"]
            bson_img = item["imgs"][img_idx]["picture"]

            # Load the image.
            img = load_img(io.BytesIO(bson_img), target_size=self.target_size)

            # Preprocess the image.
            x = img_to_array(img)
            x = self.image_data_generator.random_transform(x)
            x = self.image_data_generator.standardize(x)

            # Add the image and the label to the batch (one-hot encoded).
            batch_x[i] = x
            if self.with_labels:
                batch_y[i, image_row["category_idx"]] = 1

    if self.with_labels:
        return batch_x, batch_y
    else:
        return batch_x

def next(self):
    with self.lock:

```

```

        index_array = next(self.index_generator)
    return self._get_batches_of_transformed_samples(index_array[0])

```

```
train_bson_file = open(train_bson_path, "rb")
```

Because the training and validation generators read from the same BSON file, they need to use the same lock to protect it.

```

import threading
lock = threading.Lock()

```

Create a generator for training and a generator for validation.

```

num_classes = 5270
num_train_images = len(train_images_df)
num_val_images = len(val_images_df)
batch_size = 300

# Tip: use ImageDataGenerator for data augmentation and preprocessing.
train_datagen = ImageDataGenerator(rescale=1./255)
train_gen = BSONIterator(train_bson_file, train_images_df, train_offsets_df,
                        num_classes, train_datagen, lock,
                        batch_size=batch_size, shuffle=True)

val_datagen = ImageDataGenerator(rescale=1./255)
val_gen = BSONIterator(train_bson_file, val_images_df, train_offsets_df,
                      num_classes, val_datagen, lock,
                      batch_size=batch_size, shuffle=True)

```

```
    Found 11139021 images belonging to 5270 classes.
```

```
    Found 1232272 images belonging to 5270 classes.
```

```
next(train_gen) # warm-up
```

```
%time bx, by = next(train_gen)
```

```
    CPU times: user 416 ms, sys: 38.3 ms, total: 454 ms
```

```
    Wall time: 1.07 s
```

```
#Found 989740 images belonging to 5270 classes.
```

```
#Found 242058 images belonging to 5270 classes.
```

```
# prepare data augmentation configuration
```

```
#train_datagen = ImageDataGenerator(
```


```
#         rescale=1./255,
#         shear_range=0.2,
#         zoom_range=0.2,
#         horizontal_flip=True)
```

How fast is the generator? Create a single batch:

Does it really output images and one-hot encoded class labels? Note that the images are pre-processed (and augmented) and therefore may look weird.

```
plt.imshow(bx[-1].astype(np.uint8))
```

<matplotlib.image.AxesImage at 0x7f74e4aed9d0>



```
cat_idx = np.argmax(by[-1])
cat_id = idx2cat[cat_idx]
categories_df.loc[cat_id]
```


```
category_level1      LIBRAIRIE
category_level2      BRICOLAGE - DECORATION - JARDINAGE
category_level3      BRICOLAGE
category_idx          3259
Name: 1000014058, dtype: object
```

```
%time bx, by = next(val_gen)
```

```
CPU times: user 429 ms, sys: 90 ms, total: 519 ms
Wall time: 1.03 s
```

```
plt.imshow(bx[-1].astype(np.uint8))
```

<matplotlib.image.AxesImage at 0x7f74e45dfd90>



```
cat_idx = np.argmax(by[-1])
cat_id = idx2cat[cat_idx]
categories_df.loc[cat_id]
```

```
category_level1      SPORT
category_level2      RUNNING - ATHLETISME
category_level3      CHAUSSURES DE RUNNING - CHAUSSURES D'ATHLETISME
category_idx          4798
Name: 1000018838, dtype: object
```

▼ Part 3: Training

Create a very simple Keras model and train it, to test that the generators work.

```
from keras.models import Sequential
from keras.layers import Dropout, Flatten, Dense
from keras.layers.convolutional import Conv2D
from keras.layers.pooling import MaxPooling2D, GlobalAveragePooling2D

model = Sequential()
model.add(Conv2D(32, 3, padding="same", activation="relu", input_shape=(180, 180, 3)))
model.add(MaxPooling2D())
model.add(Conv2D(64, 3, padding="same", activation="relu"))
model.add(MaxPooling2D())
model.add(Conv2D(128, 3, padding="same", activation="relu"))
model.add(MaxPooling2D(2,2))
model.add(GlobalAveragePooling2D())
model.add(Dense(num_classes, activation="softmax"))

model.compile(optimizer="adam",
              loss="categorical_crossentropy",
              metrics=["accuracy"])

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 180, 180, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 90, 90, 32)	0
conv2d_5 (Conv2D)	(None, 90, 90, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 45, 45, 64)	0
conv2d_6 (Conv2D)	(None, 45, 45, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 22, 22, 128)	0
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 128)	0
dense_2 (Dense)	(None, 5270)	679830
Total params: 773,078		
Trainable params: 773,078		
Non-trainable params: 0		

To train the model:

```
model.fit_generator(train_gen,
                    steps_per_epoch = 50,    #num_train_images // batch_size,
                    epochs = 1,
                    validation_data = val_gen,
                    validation_steps = 10,    #num_val_images // batch_size,
                    workers = 8)
```

Epoch 1/1

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

50/50 [=====] - 564s - loss: 7.4456 - acc: 0.0059 - val_loss: 1
<keras.callbacks.History at 0x7f07a0971790>



```
model.save('/content/model',include_optimizer=True)
print("Model Saved")
```

Model Saved

```
from keras.models import load_model
demo = load_model('/content/model')
```

```
# To train the model:
```

```
demo.fit_generator(train_gen,
                    steps_per_epoch = 50,    #num_train_images // batch_size,
                    epochs = 3,
                    validation_data = val_gen,
                    validation_steps = 10,    #num_val_images // batch_size,
                    workers = 8)
```

```
Epoch 1/3
```

```
50/50 [=====] - 363s - loss: 6.9756 - acc: 0.0076 - val_loss: 1
```

```
Epoch 2/3
```

```
50/50 [=====] - 323s - loss: 6.8859 - acc: 0.0101 - val_loss: 1
```

```
Epoch 3/3
```

```
50/50 [=====] - 342s - loss: 6.8476 - acc: 0.0113 - val_loss: 1
```

```
<keras.callbacks.History at 0x7f079e6b78d0>
```



```
model.save('/content/model2',include_optimizer=True)
print("Model Saved")
```

```
Model Saved
```

```
from keras.models import load_model
demo2 = load_model('/content/model2')
```

```
# To train the model:
```

```
demo2.fit_generator(train_gen,
                    steps_per_epoch = 50,    #num_train_images // batch_size,
                    epochs = 3,
                    validation_data = val_gen,
                    validation_steps = 10,    #num_val_images // batch_size,
                    workers = 8)
```



```
Epoch 1/3
29/50 [=====>.....] - ETA: 138s - loss: 7.6707 - acc: 0.0054
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
```

```
# To evaluate on the validation set:
```

```
#model.evaluate_generator(val_gen, steps=num_val_images // batch_size, workers=8)
```

```
-----> 7                                workers = 8)
```

```
def lr_scheduler(epoch, lr):
```

```
    'For every 2nd epoch, decay learning rate by 5%'
```

```
    decay_rate = .95
```

```
    decay_step = 1
```

```
    if (epoch+1) % decay_step == 0 :
```

```
        return lr * decay_rate
```

```
    return lr
```

```
from keras.callbacks import ReduceLROnPlateau
```

```
from keras.callbacks import LearningRateScheduler
```

```
reduce_lr = ReduceLROnPlateau(monitor='val_acc', factor=0.2, patience=1, min_lr=0.000000003)
```

```
#lrschedule = LearningRateScheduler(lr_scheduler)
```

```
from keras.layers import Input, Lambda, Dense, Flatten
```

```
from keras.models import Model
```

```
from keras.applications.vgg16 import VGG16
```

```
from keras.applications.vgg16 import preprocess_input
```

```
from keras.preprocessing import image
```

```
from keras.layers import Dense, Conv2D, MaxPool2D , Flatten
```

```
from keras.callbacks import Callback
```

```
from keras.callbacks import TensorBoard
```

```
from keras.applications.xception import Xception
```

```
from keras.applications.xception import preprocess_input
```

```
IMAGE_SIZE = [180, 180] #pre trained Xception model
```

```
model = Xception(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download
```

```
79699968/83683744 [=====>..] - ETA: 0sWARNING:tensorflow:From /usr/
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
#model_1
for layer in model.layers:
    layer.trainable = False
#Adding custom Layers
x = model.output
x = Conv2D(filters=512,kernel_size=(3,3),padding="same", activation="relu")(x)
x = MaxPool2D(2,2)(x)
x = Flatten()(x)
x = Dense(500, activation="relu")(x)
x = Dense(500, activation="relu")(x)
output = Dense(5270, activation="softmax")(x)
# creating the final model
model_1 = Model(inputs = model.input, outputs = output)
# compile the model
Adam = keras.optimizers.Adam(lr=0.0003, beta_1=0.9, beta_2=0.999, epsilon=1e-07)
model_1.compile(loss = "categorical_crossentropy", optimizer = Adam, metrics=["accuracy"])

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
model_1.summary()
```

block5_sepconv1 (SeparableConv2D (None, 11, 11, 728))	536536	block5_sepconv1_activation_1
block5_sepconv1_bn (BatchNormalizer (None, 11, 11, 728))	2912	block5_sepconv1[0]
block5_sepconv2_act (Activation) (None, 11, 11, 728)	0	block5_sepconv1_bn
block5_sepconv2 (SeparableConv2D (None, 11, 11, 728))	536536	block5_sepconv2_activation_1
block5_sepconv2_bn (BatchNormalizer (None, 11, 11, 728))	2912	block5_sepconv2[0]

block5_sepconv3_act (Activation)	(None, 11, 11, 728)	0	block5_sepconv2_bn
block5_sepconv3 (SeparableConv2D)	(None, 11, 11, 728)	536536	block5_sepconv3_act
block5_sepconv3_bn (BatchNormali	(None, 11, 11, 728)	2912	block5_sepconv3[0]
add_4 (Add)	(None, 11, 11, 728)	0	block5_sepconv3_bn add_3[0][0]
block6_sepconv1_act (Activation)	(None, 11, 11, 728)	0	add_4[0][0]
block6_sepconv1 (SeparableConv2D)	(None, 11, 11, 728)	536536	block6_sepconv1_act
block6_sepconv1_bn (BatchNormali	(None, 11, 11, 728)	2912	block6_sepconv1[0]
block6_sepconv2_act (Activation)	(None, 11, 11, 728)	0	block6_sepconv1_bn
block6_sepconv2 (SeparableConv2D)	(None, 11, 11, 728)	536536	block6_sepconv2_act
block6_sepconv2_bn (BatchNormali	(None, 11, 11, 728)	2912	block6_sepconv2[0]
block6_sepconv3_act (Activation)	(None, 11, 11, 728)	0	block6_sepconv2_bn
block6_sepconv3 (SeparableConv2D)	(None, 11, 11, 728)	536536	block6_sepconv3_act
block6_sepconv3_bn (BatchNormali	(None, 11, 11, 728)	2912	block6_sepconv3[0]
add_5 (Add)	(None, 11, 11, 728)	0	block6_sepconv3_bn add_4[0][0]
block7_sepconv1_act (Activation)	(None, 11, 11, 728)	0	add_5[0][0]
block7_sepconv1 (SeparableConv2D)	(None, 11, 11, 728)	536536	block7_sepconv1_act
block7_sepconv1_bn (BatchNormali	(None, 11, 11, 728)	2912	block7_sepconv1[0]
block7_sepconv2_act (Activation)	(None, 11, 11, 728)	0	block7_sepconv1_bn
block7_sepconv2 (SeparableConv2D)	(None, 11, 11, 728)	536536	block7_sepconv2_act
block7_sepconv2_bn (BatchNormali	(None, 11, 11, 728)	2912	block7_sepconv2[0]
block7_sepconv3_act (Activation)	(None, 11, 11, 728)	0	block7_sepconv2_bn
block7_sepconv3 (SeparableConv2D)	(None, 11, 11, 728)	536536	block7_sepconv3_act
block7_sepconv3_bn (BatchNormali	(None, 11, 11, 728)	2912	block7_sepconv3[0]

To train the model:

```
model_1.fit_generator(train_gen,
                      steps_per_epoch = num_train_images // batch_size, #num_train_images //
                      epochs = 1,
                      validation_data = val_gen,
```

```
validation_steps = num_val_images // batch_size, #num_val_images // batch_size
workers = 8)
```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/tensorflow_core/python/ops/stack_ops.py:114: Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:


WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

Epoch 1/1

37130/37130 [=====] - 27517s - loss: 2.9607 - acc: 0.4601 - val_loss: 2.9607 - val_acc: 0.4601
<keras.callbacks.History at 0x7f596698aa10>



```
model_1.save('/content/drive/MyDrive/kerasep1_model.keras', include_optimizer=True)
print("Model Saved")
```

Model Saved

```
model_1.save_weights('/content/drive/MyDrive/kerasmodel1_weights.h5')
print("Weights Saved")
```

Weights Saved

```
from keras.models import load_model
demo2 = load_model('/content/drive/MyDrive/kerasep1_model.keras')
```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

keep_dims is deprecated, use keepdims instead

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_backend.py:1445: Instructions for updating:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.


```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/tensorflow_core/python/op
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the construct
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

```



```
keras.backend.set_value(demo2.optimizer.lr, 0.0003)
```

```
demo2.optimizer.get_config()
```

```
{'lr': 0.0003000000142492354,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
# To train the model:
```


```
demo2.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size, #num_train_images //
                    epochs = 2,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size, #num_val_images // batc
                    workers = 8,
                    callbacks= [reduce_lr])
```

```
Epoch 1/2
```

```
37130/37130 [=====] - 27757s - loss: 2.5840 - acc: 0.5098 - val
(self.monitor, ','.join(list(logs.keys()))), RuntimeWarning
```

```
Epoch 2/2
```

```
37130/37130 [=====] - 27764s - loss: 2.4930 - acc: 0.5228 - val
<keras.callbacks.History at 0x7fdb7257e910>
```



```
demo2.save('/content/drive/MyDrive/kerasep3_model.keras', include_optimizer=True)
print("Model Saved")
```

```
Model Saved
```

```
model_1.save_weights('/content/drive/MyDrive/kerasmodelp3_weight')
print("Wights Saved")
```

```
from keras.models import load_model
demo2 = load_model('/content/drive/MyDrive/kerasep3_model.kerasave')
```

To train the model:

```
demo2.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size,    #num_train_images //
                    epochs = 2,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size,    #num_val_images // batc
                    workers = 8,
                    callbacks= [reduce_lr])
```

Epoch 1/2

37130/37130 [=====] - 27658s - loss: 2.4442 - acc: 0.5298 - val

Epoch 2/2

37130/37130 [=====] - 27675s - loss: 2.4123 - acc: 0.5344 - val

<keras.callbacks.History at 0x7f9ff1f87510>



```
demo2.save('/content/drive/MyDrive/kerasep4_model.kerasave',include_optimizer=True)
print("Model Saved")
```

Model Saved

Double-click (or enter) to edit

```
from keras.models import load_model
demo2 = load_model('/content/drive/MyDrive/kerasep5_model.kerasave')
```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

Instructions for updating:

keep_dims is deprecated, use keepdims instead

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_


WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/tensorflow_core/python/op
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the construct
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

```



```
demo2.optimizer.get_config()
```

```
{'lr': 0.0003000000142492354,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
0.0003*0.65
```

```
0.000195
```

```
keras.backend.set_value(demo2.optimizer.lr, 0.000195)
```

```
demo2.optimizer.get_config()
```

```
{'lr': 0.00019500000053085387,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
# To train the model:
```

```
demo2.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size,    #num_train_images //
                    epochs = 3,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size,    #num_val_images // batc
                    workers = 8,
                    callbacks= [reduce_lr])
```

Epoch 1/3


```
37130/37130 [=====] - 27764s - loss: 2.2950 - acc: 0.5504 - val
(self.monitor, ', '.join(list(logs.keys()))), RuntimeWarning
```

Epoch 2/3

```
37130/37130 [=====] - 27759s - loss: 2.2476 - acc: 0.5567 - val
```

Epoch 3/3

```
37130/37130 [=====] - 27760s - loss: 2.2218 - acc: 0.5600 - val
<keras.callbacks.History at 0x7f75a02644d0>
```



```
demo2.save('/content/drive/MyDrive/kerasep8_model.keras', include_optimizer=True)
print("Model Saved")
```

Model Saved

▼ lr 0.00005

```
from keras.models import load_model
demo2 = load_model('/content/drive/MyDrive/kerasep8_model.keras')
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
```

```
keep_dims is deprecated, use keepdims instead
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```



```


WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711:

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/tensorflow_core/python/op
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the construct
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

```



```
demo2.optimizer.get_config()
```

```
{'lr': 0.00019500000053085387,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
keras.backend.set_value(demo2.optimizer.lr, 0.00005)
```

```
demo2.optimizer.get_config()
```

```
{'lr': 4.999999873689376e-05,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
# To train the model:
```

```
demo2.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size, #num_train_images //
                    epochs = 3,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size, #num_val_images // batc
                    workers = 8,
                    callbacks= [reduce_lr])
```

```
Epoch 1/3
```

```
37130/37130 [=====] - 27458s - loss: 2.0836 - acc: 0.5791 - val
```

```
Epoch 2/3
```

```
37130/37130 [=====] - 27392s - loss: 2.0411 - acc: 0.5849 - val
```

```
Epoch 3/3
37130/37130 [=====] - 27400s - loss: 2.0189 - acc: 0.5879 - val
<keras.callbacks.History at 0x7f8ef5362d90>
```



```
demo2.save('/content/drive/MyDrive/kerasep11_model.kerasave', include_optimizer=True)
print("Model Saved")
```

Model Saved

Double-click (or enter) to edit

```
from keras.models import load_model
demo2 = load_model('/content/drive/MyDrive/kerasep11_model.kerasave')


demo2.optimizer.get_config()

keras.backend.set_value(demo2.optimizer.lr, 0.000005)

demo2.optimizer.get_config()

# To train the model:
demo2.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size, #num_train_images //
                    epochs = 1,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size, #num_val_images // batc
                    workers = 8,
                    callbacks= [reduce_lr])
```

```
Epoch 1/1
37130/37130 [=====] - 27535s - loss: 1.9701 - acc: 0.5944 - val
<keras.callbacks.History at 0x7fc64a02ac10>
```



```
demo2.save('/content/drive/MyDrive/kerasep12_model.kerasave', include_optimizer=True)
print("Model Saved")
```

Model Saved

▼ n

```
from keras.models import load_model
```

```
model = load_model('/content/drive/MyDrive/kerasep12_model.kerasave')
```

```
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711:
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the construct
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
```

```
model.summary()
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 180, 180, 3)	0	
block1_conv1 (Conv2D)	(None, 89, 89, 32)	864	input_1[0][0]
block1_conv1_bn (BatchNormalizat	(None, 89, 89, 32)	128	block1_conv1[0][0]
block1_conv1_act (Activation)	(None, 89, 89, 32)	0	block1_conv1_bn[0]
block1_conv2 (Conv2D)	(None, 87, 87, 64)	18432	block1_conv1_act[0]
block1_conv2_bn (BatchNormalizat	(None, 87, 87, 64)	256	block1_conv2[0][0]
block1_conv2_act (Activation)	(None, 87, 87, 64)	0	block1_conv2_bn[0]
block2_sepconv1 (SeparableConv2D	(None, 87, 87, 128)	8768	block1_conv2_act[0]
block2_sepconv1_bn (BatchNormali	(None, 87, 87, 128)	512	block2_sepconv1[0]
block2_sepconv2_act (Activation)	(None, 87, 87, 128)	0	block2_sepconv1_bn
block2_sepconv2 (SeparableConv2D	(None, 87, 87, 128)	17536	block2_sepconv2_ac
block2_sepconv2_bn (BatchNormali	(None, 87, 87, 128)	512	block2_sepconv2[0]
conv2d_1 (Conv2D)	(None, 44, 44, 128)	8192	block1_conv2_act[0]
block2_pool (MaxPooling2D)	(None, 44, 44, 128)	0	block2_sepconv2_bn
batch_normalization_1 (BatchNorm	(None, 44, 44, 128)	512	conv2d_1[0][0]

add_1 (Add)	(None, 44, 44, 128)	0	block2_pool[0][0] batch_normalizatio
block3_sepconv1_act (Activation)	(None, 44, 44, 128)	0	add_1[0][0]
block3_sepconv1 (SeparableConv2D)	(None, 44, 44, 256)	33920	block3_sepconv1_ac
block3_sepconv1_bn (BatchNormali	(None, 44, 44, 256)	1024	block3_sepconv1[0]
block3_sepconv2_act (Activation)	(None, 44, 44, 256)	0	block3_sepconv1_bn
block3_sepconv2 (SeparableConv2D)	(None, 44, 44, 256)	67840	block3_sepconv2_ac
block3_sepconv2_bn (BatchNormali	(None, 44, 44, 256)	1024	block3_sepconv2[0]
conv2d_2 (Conv2D)	(None, 22, 22, 256)	32768	add_1[0][0]
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0	block3_sepconv2_bn
batch_normalization_2 (BatchNorm	(None, 22, 22, 256)	1024	conv2d_2[0][0]
add_2 (Add)	(None, 22, 22, 256)	0	block3_pool[0][0] batch_normalizatio

▼ set last 3 layer of xception net to trainable

```
my_layer = model.get_layer('block14_sepconv2')
my_layer.trainable = True
```

```
my_layer = model.get_layer('block14_sepconv2_bn')
my_layer.trainable = True
```

```
my_layer = model.get_layer('block14_sepconv2_act')
my_layer.trainable = True
```

```
model.summary()
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	(None, 180, 180, 3)	0	
block1_conv1 (Conv2D)	(None, 89, 89, 32)	864	input_1[0][0]
block1_conv1_bn (BatchNormalizat	(None, 89, 89, 32)	128	block1_conv1[0][0]
block1_conv1_act (Activation)	(None, 89, 89, 32)	0	block1_conv1_bn[0]

block1_conv2 (Conv2D)	(None, 87, 87, 64)	18432	block1_conv1_act[0]
block1_conv2_bn (BatchNormalizat	(None, 87, 87, 64)	256	block1_conv2[0][0]
block1_conv2_act (Activation)	(None, 87, 87, 64)	0	block1_conv2_bn[0]
block2_sepconv1 (SeparableConv2D	(None, 87, 87, 128)	8768	block1_conv2_act[0]
block2_sepconv1_bn (BatchNormali	(None, 87, 87, 128)	512	block2_sepconv1[0]
block2_sepconv2_act (Activation)	(None, 87, 87, 128)	0	block2_sepconv1_bn
block2_sepconv2 (SeparableConv2D	(None, 87, 87, 128)	17536	block2_sepconv2_ac
block2_sepconv2_bn (BatchNormali	(None, 87, 87, 128)	512	block2_sepconv2[0]
conv2d_1 (Conv2D)	(None, 44, 44, 128)	8192	block1_conv2_act[0]
block2_pool (MaxPooling2D)	(None, 44, 44, 128)	0	block2_sepconv2_bn
batch_normalization_1 (BatchNorm	(None, 44, 44, 128)	512	conv2d_1[0][0]
add_1 (Add)	(None, 44, 44, 128)	0	block2_pool[0][0] batch_normalizatio
block3_sepconv1_act (Activation)	(None, 44, 44, 128)	0	add_1[0][0]
block3_sepconv1 (SeparableConv2D	(None, 44, 44, 256)	33920	block3_sepconv1_ac
block3_sepconv1_bn (BatchNormali	(None, 44, 44, 256)	1024	block3_sepconv1[0]
block3_sepconv2_act (Activation)	(None, 44, 44, 256)	0	block3_sepconv1_bn
block3_sepconv2 (SeparableConv2D	(None, 44, 44, 256)	67840	block3_sepconv2_ac
block3_sepconv2_bn (BatchNormali	(None, 44, 44, 256)	1024	block3_sepconv2[0]
conv2d_2 (Conv2D)	(None, 22, 22, 256)	32768	add_1[0][0]
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0	block3_sepconv2_bn
batch_normalization_2 (BatchNorm	(None, 22, 22, 256)	1024	conv2d_2[0][0]
add_2 (Add)	(None, 22, 22, 256)	0	block3_pool[0][0] batch_normalizatio

model.optimizer.get_config()

```
{'lr': 4.999999873689376e-06,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
keras.backend.set_value(model.optimizer.lr, 0.0001)
```

```
model.optimizer.get_config()
```

```
{'lr': 9.999999747378752e-05,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
# To train the model:
```

```
model.fit_generator(train_gen,
                    steps_per_epoch = num_train_images // batch_size,  #num_train_images //
                    epochs = 2,
                    validation_data = val_gen,
                    validation_steps = num_val_images // batch_size,  #num_val_images // batch_size
                    workers = 8,
                    callbacks= [reduce_lr])
```

```
Epoch 1/2
37130/37130 [=====] - 27447s - loss: 2.0486 - acc: 0.5834 - val
Epoch 2/2
37130/37130 [=====] - 27441s - loss: 2.0400 - acc: 0.5845 - val
<keras.callbacks.History at 0x7fa0d7333b90>
```

```
model.save('/content/drive/MyDrive/kerasep14_model.keras', include_optimizer=True)
print("Model Saved")
```

Model Saved

Double-click (or enter) to edit

```
from keras.models import load_model
model2 = load_model('/content/drive/MyDrive/kerasep14_model.keras')
```

```
/usr/local/lib/python3.7/dist-packages/keras/models.py:287: UserWarning: Error in loading
warnings.warn('Error in loading the saved optimizer ')
```

```
model2.optimizer.get_config()
```

```
{'lr': 9.999999747378752e-05,
 'beta_1': 0.8999999761581421,
 'beta_2': 0.9990000128746033,
 'decay': 0.0,
 'epsilon': 1e-07}
```

```
keras.backend.set_value(model.optimizer.lr, 0.00005)
```

▼ Prediction

```
!kaggle competitions download -c 'cdiscout-image-classification-challenge' -f test.bson
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you
Downloading test.bson.zip to /content
100% 11.8G/11.8G [02:33<00:00, 74.5MB/s]
100% 11.8G/11.8G [02:33<00:00, 82.5MB/s]
```




```
!unzip /content/test.bson.zip
```

```
Archive: /content/test.bson.zip
  inflating: test.bson
```

```
!rm /content/test.bson.zip
```

```
!kaggle competitions download -c 'cdiscout-image-classification-challenge' -f sample_submiss
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you
Downloading sample_submission.csv.zip to /content
100% 4.61M/4.61M [00:00<00:00, 28.5MB/s]
100% 4.61M/4.61M [00:00<00:00, 28.4MB/s]
```



```
!unzip /content/sample_submission.csv.zip
```

```
Archive: /content/sample_submission.csv.zip
  inflating: sample_submission.csv
```


```
from keras.models import load_model
model = load_model('/content/drive/MyDrive/kerasep12_model.kerasave')
```

```
Using TensorFlow backend.
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
```

```

WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/optimizers.py:711:
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
keep_dims is deprecated, use keepdims instead
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/tensorflow_core/python/op
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to the construct
WARNING:tensorflow:From /usr/local/lib/python3.7/dist-packages/keras/backend/tensorflow_

```



```

from keras import backend as K
from keras.preprocessing.image import ImageDataGenerator
from keras.models import load_model
import pandas as pd
import bson
import os
from collections import defaultdict
from tqdm import *
import multiprocessing as mp
import struct
import os, sys, math, io
import numpy as np
from keras.preprocessing.image import load_img, img_to_array

num_test_products = 1768182

categories_path = os.path.join("/content/drive/Shared drives/Case_study/cdiscount/category_nam
categories_df = pd.read_csv(categories_path, index_col="category_id")

# Maps the category_id to an integer index. This is what we'll use to
# one-hot encode the labels.
categories_df["category_idx"] = pd.Series(range(len(categories_df)), index=categories_df.inde

```



```

categories_df.to_csv("categories.csv")
categories_df.head()

def make_category_tables():
    cat2idx = {}
    idx2cat = {}
    for ir in categories_df.itertuples():
        category_id = ir[0]
        category_idx = ir[4]
        cat2idx[category_id] = category_idx
        idx2cat[category_idx] = category_id
    return cat2idx, idx2cat

cat2idx, idx2cat = make_category_tables()

submission_df = pd.read_csv("/content/sample_submission.csv")
submission_df.head()

test_bson_path = os.path.join("/content/test.bson")

# test_datagen = ImageDataGenerator(
#     rescale=1./255,
#     shear_range=0.2,
#     zoom_range=0.3,
#     rotation_range=180.,
#     width_shift_range=0.3,
#     height_shift_range=0.3,
#     horizontal_flip=True)

test_datagen = ImageDataGenerator(
    rescale=1./255)

data = bson.decode_file_iter(open(test_bson_path, "rb"))

with tqdm(total=num_test_products) as pbar:
    for c, d in enumerate(data):
        product_id = d["_id"]
        num_imgs = len(d["imgs"])
        batch_x = np.zeros((num_imgs, 180, 180, 3), dtype=K.floatx())
        for i in range(num_imgs):
            bson_img = d["imgs"][i]["picture"]
            # Load and preprocess the image.
            img = load_img(io.BytesIO(bson_img), target_size=(180, 180))
            x = img_to_array(img)
            #x = test_datagen.random_transform(x)
            x = test_datagen.standardize(x)
            # Add the image to the batch.
            batch_x[i] = x
        prediction = model.predict(batch_x, batch_size=num_imgs)
        avg_pred = prediction.mean(axis=0) # .mean
        cat_idx = np.argmax(avg_pred)
        submission_df.iloc[c]["category_id"] = idx2cat[cat_idx]

```

```
pbar.update()

submission_df.to_csv("final_submission.csv.gz", compression="gzip", index=False)

100%|██████████| 1768182/1768182 [7:40:13<00:00, 64.03it/s]

import shutil

# Source path 100006949
source = "/content/final_submission.csv.gz"

# Destination path
destination = "/content/drive/MyDrive"

# Move the content of
# source to destination
dest = shutil.move(source, destination)
print('done')

done
```

[Colab paid products](#) - [Cancel contracts here](#)

7m 15s completed at 11:10 PM

